

# Object Oriented Programming - Exercise 4: AVL Trees

## 1 Goals

1. Implementing the AVL Tree data structure.
2. Experimenting with this data structure.
3. Write a design in the UML format.

## 2 Submission Details

- Code submission Deadline: **Wednesday, 17/05/2017, 23:55**
- UML submission Deadline: **Wednesday, 10/05/2017, 23:55**
- This exercise will be done alone (**not in pairs**).

## 3 Java Classes & Interfaces

You may not use any classes and interfaces you didn't write yourself, except the following:

- `java.lang.Iterable<T>`
- `java.util.Iterator<T>`
- `java.util.List<T>`
- `java.util.LinkedList<T>`
- `java.util.ArrayList<T>`
- `Math`
- The class `Exception` and the classes that extends it

## 4 Introduction

A binary search tree is a tree in which all nodes have at most two children, and each node's value is smaller than all the values in his right subtree and bigger than all the values in his left subtree. An AVL tree is a self-balancing binary search tree. AVL trees maintain the *AVL property* - for each node, the difference between the heights of both of its sub-trees is at most 1.<sup>1</sup> This is done by re-balancing the tree after each insertion and deletion operation. You've recently learned about the theoretical background of AVL trees in the Data Structures (DaSt) course.<sup>2</sup>

In this exercise, you will implement an AVL tree of integers based on the pseudo-code shown in the DaSt lecture and tirgul. Duplicate keys will not be allowed in the tree.

## 5 API

You are required to submit a class named `AvlTree` that implements `Iterable<Integer>` and the following API.

You are allowed (and encouraged!) to write and submit more class(es) as you see fit. You can also add methods of your own but make sure that you don't change the supplied API; you are only allowed to add methods with the `private` and the `default-package` access modifier.

### 5.1 Package

Both `AvlTree` and any other class you implement in this exercise should be placed in the `oop.ex4.data_structures` package, since you build a data-structure that will be used by other users and has no entry-point (main) of its own.

### 5.2 Methods

- ```
/**
 * Add a new node with the given key to the tree.
 *
 * @param newValue the value of the new node to add.
 * @return true if the value to add is not already in the tree and it was successfully added,
 * false otherwise.
 */
public boolean add(int newValue);
```
- ```
/**
 * Check whether the tree contains the given input value.
 *
 * @param searchVal the value to search for.
 * @return the depth of the node (0 for the root) with the given value if it was found in
 * the tree, -1 otherwise.
 */
public int contains(int searchVal);
```

---

<sup>1</sup>The height of a tree is defined as the length of the longest downward path from the root to any of the leaves.

<sup>2</sup>See <http://moodle.cs.huji.ac.il/cs13/file.php/67109/tirgul8.pdf>.

- ```
/**
 * Removes the node with the given value from the tree, if it exists.
 *
 * @param toDelete the value to remove from the tree.
 * @return true if the given value was found and deleted, false otherwise.
 */
public boolean delete(int toDelete);
```
- ```
/**
 * @return the number of nodes in the tree.
 */
public int size();
```
- ```
/**
 * @return an iterator for the AVL Tree. The returned iterator iterates over the tree nodes
 * in an ascending order, and does NOT implement the remove() method.
 */
public Iterator<Integer> iterator();
```

### 5.3 Constructors

- ```
/**
 * The default constructor.
 */
public AVLTree();
```
- ```
/**
 * A constructor that builds a new AVL tree containing all unique values in the input
 * array.
 * @param data the values to add to tree.
 */
public AVLTree(int[] data);
```
- ```
/**
 * A copy constructor that creates a deep copy of the given AVLTree. The new tree
 * contains all the values of the given tree, but not necessarily in the same structure.
 * @param avlTree an AVL tree.
 */
public AVLTree(AVLTree avlTree);
```

**Comment:** A deep copy of the tree means that for every node or any other internal object of the given tree, a new, identical object, is instantiated for the new tree (the internal object is not simply referenced from it; i.e. points to the same internal node object in memory). You are not required to implement this method with the best possible worst-case complexity.

## 5.4 Static Methods

**Comment:** You are required to implement these methods with a running time complexity that is at most polynomial in  $h$ , and uses  $O(1)$  memory.

- ```
/**
 * Calculates the minimum number of nodes in an AVL tree of height h.
 *
 * @param h the height of the tree (a non-negative number) in question.
 * @return the minimum number of nodes in an AVL tree of the given height.
 */
public static int findMinNodes(int h);
```
- ```
/**
 * Calculates the maximum number of nodes in an AVL tree of height h.
 *
 * @param h the height of the tree (a non-negative number) in question.
 * @return the maximum number of nodes in an AVL tree of the given height.
 */
public static int findMaxNodes(int h);
```

For example, `findMinNodes(3)`, as shown in Figure 1, should return 7. `findMaxNodes(3)` should return 15.

## 5.5 Comments

1. There are many types of binary search trees: red-black trees, splay trees, treap, T-trees and more. A good programmer (that likes high grades) should think about these type of things when planning a design, even if it is not currently demanded in the project. On the other hand, avoid extreme over-engineering and other [anti-patterns](#). A design that can be easily changed to support future reasonable requirements will get better grades.
2. The `AvlTree(int[] data)` constructor should construct a tree by creating an empty one and then adding the elements in the input array one by one. Note that if a value appears more than once in the list, only the first appearance is added.
3. The `add(int newValue)` method will return false and leave the tree unchanged if `newValue` already exists in it.
4. The data constructor can receive any array of `int` primitives, which means **you can't assume the array is sorted**. You also can't assume that you won't get a null reference.
5. You can assume `findMinNodes()` and `findMaxNodes()` receive a non-negative integer.
6. You are expected to throw exceptions when needed.
7. **Do not** paste code copied from the PDF file. It contains unwanted characters that will fail compilation.

8. A tree with only a single node is of height zero, or has  $h = 0$ . A tree with a root that has one or two sons (each of which has no sons) is of height one, or  $h = 1$ . Etc.
9. When writing a class implementing an interface only partially, i.e., not implementing some optional methods of the interface, you should throw an `UnsupportedOperationException` from the methods you wish to leave unimplemented. Can you figure out how throwing this specific exception does not violate the interface's API and method signature?
10. The return type of the `iterator()` method must be `Iterator<Integer>`, and not an iterator of any other type.
11. You may have `AvlTree` extend another class, but then you may not add new `public` or `protected` methods to this base class (as `AvlTree` will inherit any such method). In such a case, you may have the base class implement `Iterable<Integer>` instead of `AvlTree` itself.
12. You may also have `AvlTree` implement other interfaces. Make sure, though, that this does not entail extending the API, and that you can explain your design choices.

## 6 Theoretical Questions


Answer the following questions, and write your full answers in your README file:

### 6.1 Analyzing the AVL tree

In Figure 1 you can see an AVL tree of height 3. This tree may seem unbalanced, but is in fact a valid AVL tree resulted from a specific insertion order. Notice that this example shows a tree of height 3 with the minimal number of nodes (if you delete any node, the height of the tree will become 2).

1. Find a series of 12 numbers, such that when they are inserted into an empty AVL tree one-by-one, the result will be a tree of height 4 (insertions only, no deletions).

Hint: We suggest you start by figuring out the order in which the nodes were inserted into the example tree, and then continue to create the larger one.



**Figure 1:** an AVL tree of height 3 with the minimal number of nodes

## 6.2 The complexity of constructing an AVL tree

1. What is the asymptotic running time complexity of your tree construction from an array of length  $n$  done in `AvlTree(int[] data)`? Explain briefly.
2. What is the asymptotic running time complexity of your tree construction done in `AvlTree(AvlTree avlTree)`? Explain briefly.

## 7 UML Submission Requirements

A week before your code submission you will submit a single PDF file that contains your design in a UML diagram. You can scan a readable hand written one or generate the diagram with IntelliJ [built-in framework](#). This will make you stop and plan before jumping into code. You can add comments under the diagram that explain your decisions. Although you are requested to submit a well planned design, you are allowed to change it if you have a good reason. In that case you are to write an explanation for your changes in the README of the code submission.

Note: In real life, big agile software projects use UML-like diagrams to exchange ideas in a clear way between people, but one shouldn't look at his design as a closed contract. Even very experienced developers find unexpected discoveries along the way that force them to change course. See Wikipedia's [Agile software development](#) page for more about modern development concepts that are common in start-ups and small software companies.

## 8 Code Submission Requirements

### 8.1 README

Please address the following points in your README file:

1. Describe which class(es) (if any) you wrote as part of your implementation of an AVL tree, other than `AvlTree`. The description should include the purpose of each class, its important methods and its interaction with the `AvlTree` class.
2. Describe your implementation of the methods `add()` and `delete()`. The description should include the general workflow in each of these methods. You should also indicate which helper methods you implemented for each of them, and which of these helper methods are shared by both of them (if any).
3. Your README file should also answer the question from Section 6.

### 8.2 Automatic Testing

Our automatic tests will do the following:

1. Make sure your code compiles and contains a README file.
2. Test your tree on interesting AVL scenarios (i.e., some of the rotation cases).<sup>3</sup>

---

<sup>3</sup>An explanation in Hebrew of the four cases can be found here: [https://he.wikipedia.org/wiki/%D7%A2%D7%A5\\_AVL](https://he.wikipedia.org/wiki/%D7%A2%D7%A5_AVL).

3. Test your code on more complicated scenarios that include large inputs and multiple add/delete operations.
4. We will run unseen tests on your code.

We will compile your class with our own tester and run a set of tests. Each such test is defined by an input file, which creates an AVL tree and performs a set of operations on it. The format of the input files is described in Appendix A.

Note: You are encouraged to create your own tests. Test all the rotation cases, complex usage, multiple insertions and deletions, etc.

## 8.3 Submission Guidelines

In the code submission, you should submit a file named `ex4.jar` containing all the `.java` files of your program, as well as the `README` file. Please note the following:

- Files should be submitted in the original directory hierarchy of their packages.
- No `.class` files should be submitted.
- There is no need to submit any testers.
- Your program must compile without any errors or warnings. IDEA displays most warnings in the editor. In order to also have them displayed when you compile your code, go to `File->Settings->Build, Execution, Deployment->Compiler->Java Compiler` and enter the following line under `Additional command line parameters`:  
`-Xlint:rawtypes -Xlint:static -Xlint:empty -Xlint:divzero -Xlint:deprecation`  
If you want to check for them using the console, you can use the following customized compilation command:  
`javac -Xlint:rawtypes -Xlint:static -Xlint:empty -Xlint:divzero -Xlint:deprecation file1.java file2.java ...`  
Note: Copy-pasting the command above to the shell might result in some invalid characters.
- `javadoc` should compile without any errors or warnings.

You may use the following unix command to create the jar files:

```
jar -cvf ex4.jar README oop/ex4/data_structures/*.java
```

This command should be run from the main project directory (that is, where the `oop` directory is).

## 9 Technical Issues

A file named `ex4_files.zip` can be found in the course website. This zip contains the test files on which your program will be tested. When you submit your exercise, we will run your code against these input files (as well as others), and you will receive a response file with the tests you

failed (if any). Error messages will contain some information about the problem. For example, you may receive the following error:

```
runTests[17](Ex4Tester):  ex4/tests/020.txt (# simple right-left rotation after delete):  
line number 7:  expected:<true> but was:<false>
```

This means that your program failed the test in file `020.txt`, when trying to execute the delete operation. More specifically, line 7 of the test failed at this stage because the return value of calling `delete()` was true instead of false (details in appendix A).

## 10 Presubmission Script

You can run the presubmission script from the shell using the following command:

```
~oop/bin/ex4/presubmission/test.py ex4.jar
```

(The same script is run when you upload your submission to Moodle.)

# Good Luck!



# Appendices

## A Test Input Files Format

Each of the automatic tests is defined in a file of the following format:

- The first line starts with a ‘#’ sign, and is a comment describing the test (i.e., this line is ignored).
- The second line creates an `AvlTree` object (denoted `myTree`) using one of the constructors defined in Section 5.3:
  - An empty line calls the default constructor.
  - A list of integers (e.g., `1 5 7 12 4`) calls the data constructor with this list of integers.
- The following lines specify a call to one of the methods of the created tree, and are executed sequentially.

They can be one of the following:

- *add number return\_value*  
call `myTree.add(number)`. Verify that the method’s return value matches *return\_value* (either `true` or `false`).
- *delete number return\_value*  
call `myTree.delete(number)`. Verify that the method’s return value matches *return\_value* (either `true` or `false`).
- *contains number return\_value*  
call `myTree.contains(number)`. Verify that the method’s return value matches *return\_value* (a non-negative integer or -1).
- *size return\_value*  
call `myTree.size()`. Verify that the method’s return value matches *return\_value* (a non-negative integer).
- *copy*  
call `myTree2 = newAvlTree(myTree)`. Verify that the copy constructor functions correctly verifying the trees contain the exact same values .
- *minNode number return\_value*  
call `AvlTree.findMinNodes(number)`. Verify that the method’s return value matches *return\_value*.

### A.1 Examples

#### A.1.1

Consider the following example test file:

```
# a simple creation of a tree and adding the number 5
```

```
add 5 true
size 1
```

In this example, the first line is a comment describing the test. The second line is empty, which means calling the empty constructor. The third line calls `myTree.add(5)`, and then verifies that the command succeeded (by verifying that the return value of the `add` method is `true`). The last line verifies that the tree size is now 1 (after adding a single element).

### A.1.2

Consider another example:

```
# calling the data constructor with the values (1 2 3) , deleting 3, and then searching for 4
1 2 3
delete 3 true
contains 4 -1
```

The first line is, again, a comment describing the test. The second line is a list of numbers, which means calling the data constructor with the integer array (1 2 3). The third line calls `myTree.delete(3)`, and then verifies that the command succeeded. The last line calls `myTree.contains(4)` and verifies that 4 is not found in the tree (return value should be -1).

## A.2 Some Clarifications

Please note the following:

- Your code passes a given test only if your methods return the correct values in **each** of the method calls.
- A test file may include numerous calls to `add`, `delete`, `contains` and `size`.